

Data 621 Assignment 1

Group2 - Gabriella Martinez, Maliat Islam, Ken Popkin, George Cruz Deschamps,

Matthew Lucich, and Karim Hammoud

2021-09-22

Table of Contents

DATA EXPLORATION:	2
Load the Data	3
DATA PREPARATION	10
BUILD MODELS	12
Model #1	12
Two predictors: Base hits by batters and Hits allowed	12
Model #2	13
Four predictors: Base hits by batters, Hits allowed, Errors, and Walks allowed	13
Model #3	14
BSR Model (SaberMetrics) (data imputation)	14
Model #4	16
(Modified) Backward Elimination Model (omitting NAs)	16
SELECT MODELS	17
Verifying OLS Regression Assumptions	17
Model Selection	20
Matt's References	20
Code Appendix	21

DATA 621 – Business Analytics and Data Mining

Homework #1 Assignment Requirements

Overview

In this homework assignment, you will explore, analyze and model a data set containing approximately 2200 records. Each record represents a professional baseball team from the years 1871 to 2006 inclusive. Each record has the performance of the team for the given year, with all of the statistics adjusted to match the performance of a 162 game season.

Your objective is to build a multiple linear regression model on the training data to predict the number of wins for the team. You can only use the variables given to you (or variables that you derive from the variables provided). Below is a short description of the variables of interest in the data set:

VARIABLE NAME	DEFINITION	THEORETICAL EFFECT
INDEX	Identification Variable (do not use)	None
TARGET_WINS	Number of wins	
TEAM_BATTING_H	Base Hits by batters (1B,2B,3B,HR)	Positive Impact on Wins
TEAM_BATTING_2B	Doubles by batters (2B)	Positive Impact on Wins
TEAM_BATTING_3B	Triples by batters (3B)	Positive Impact on Wins
TEAM_BATTING_HR	Homeruns by batters (4B)	Positive Impact on Wins
TEAM_BATTING_BB	Walks by batters	Positive Impact on Wins
TEAM_BATTING_HBP	Batters hit by pitch (get a free base)	Positive Impact on Wins
TEAM_BATTING_SO	Strikeouts by batters	Negative Impact on Wins
TEAM_BASERUN_SB	Stolen bases	Positive Impact on Wins
TEAM_BASERUN_CS	Caught stealing	Negative Impact on Wins
TEAM_FIELDING_E	Errors	Negative Impact on Wins
TEAM_FIELDING_DP	Double Plays	Positive Impact on Wins
TEAM_PITCHING_BB	Walks allowed	Negative Impact on Wins
TEAM_PITCHING_H	Hits allowed	Negative Impact on Wins
TEAM_PITCHING_HR	Homeruns allowed	Negative Impact on Wins
TEAM_PITCHING_SO	Strikeouts by pitchers	Positive Impact on Wins

Deliverables:

- A write-up submitted in PDF format. Your write-up should have four sections. Each one is described below. You may assume you are addressing me as a fellow data scientist, so do not need to shy away from technical details.
- Assigned predictions (the number of wins for the team) for the evaluation data set.
- Include your R statistical programming code in an Appendix.

In this data set we are trying to identify good and bad teams in major league baseball team's season. We are assuming some of the predictors will be higher for good teams. We will try to predict how many times a team will win in this season.

DATA EXPLORATION:

We can observe the response variable (TARGET_WINS) looks to be normally distributed. This supports the working theory that there are good teams and bad teams. There are also a lot of average teams.

There are also quite a few variables with missing values. and, Some variables are right skewed (TEAM_BASERUN_CS, TEAM_BASERUN_SB, etc.). This might support the good team theory. It may also introduce non-normally distributed residuals in the model. We shall see.

Load the Data

Summary of the train data

```
##      INDEX      TARGET_WINS      TEAM_BATTING_H TEAM_BATTING_2B
## Min.   : 1.0    Min.   : 0.00    Min.   : 891    Min.   : 69.0
## 1st Qu.: 630.8  1st Qu.: 71.00    1st Qu.:1383   1st Qu.:208.0
## Median :1270.5  Median : 82.00    Median :1454   Median :238.0
## Mean   :1268.5  Mean   : 80.79    Mean   :1469   Mean   :241.2
## 3rd Qu.:1915.5  3rd Qu.: 92.00    3rd Qu.:1537   3rd Qu.:273.0
## Max.   :2535.0  Max.   :146.00    Max.   :2554   Max.   :458.0
##
## TEAM_BATTING_3B TEAM_BATTING_HR TEAM_BATTING_BB TEAM_BATTING_SO
## Min.   : 0.00    Min.   : 0.00    Min.   : 0.0    Min.   : 0.0
## 1st Qu.: 34.00    1st Qu.: 42.00    1st Qu.:451.0    1st Qu.: 548.0
## Median : 47.00    Median :102.00    Median :512.0    Median : 750.0
## Mean   : 55.25    Mean   : 99.61    Mean   :501.6    Mean   : 735.6
## 3rd Qu.: 72.00    3rd Qu.:147.00    3rd Qu.:580.0    3rd Qu.: 930.0
## Max.   :223.00    Max.   :264.00    Max.   :878.0    Max.   :1399.0
##
##                                     NA's :102
## TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_BATTING_HBP TEAM_PITCHING_H
## Min.   : 0.0    Min.   : 0.0    Min.   :29.00    Min.   : 1137
## 1st Qu.: 66.0    1st Qu.: 38.0    1st Qu.:50.50    1st Qu.: 1419
## Median :101.0    Median : 49.0    Median :58.00    Median : 1518
## Mean   :124.8    Mean   : 52.8    Mean   :59.36    Mean   : 1779
## 3rd Qu.:156.0    3rd Qu.: 62.0    3rd Qu.:67.00    3rd Qu.: 1682
## Max.   :697.0    Max.   :201.0    Max.   :95.00    Max.   :30132
## NA's :131      NA's :772      NA's :2085
## TEAM_PITCHING_HR TEAM_PITCHING_BB TEAM_PITCHING_SO TEAM_FIELDING_E
## Min.   : 0.0    Min.   : 0.0    Min.   : 0.0    Min.   : 65.0
## 1st Qu.: 50.0    1st Qu.: 476.0    1st Qu.: 615.0    1st Qu.: 127.0
## Median :107.0    Median : 536.5    Median : 813.5    Median : 159.0
## Mean   :105.7    Mean   : 553.0    Mean   : 817.7    Mean   : 246.5
## 3rd Qu.:150.0    3rd Qu.: 611.0    3rd Qu.: 968.0    3rd Qu.: 249.2
## Max.   :343.0    Max.   :3645.0    Max.   :19278.0    Max.   :1898.0
##
##                                     NA's :102
## TEAM_FIELDING_DP
## Min.   : 52.0
## 1st Qu.:131.0
## Median :149.0
## Mean   :146.4
## 3rd Qu.:164.0
## Max.   :228.0
## NA's :286
```

##	INDEX	TEAM_BATTING_H	TEAM_BATTING_2B	TEAM_BATTING_3B
##	Min. : 9	Min. : 819	Min. : 44.0	Min. : 14.00
##	1st Qu.: 708	1st Qu.:1387	1st Qu.:210.0	1st Qu.: 35.00
##	Median :1249	Median :1455	Median :239.0	Median : 52.00
##	Mean :1264	Mean :1469	Mean :241.3	Mean : 55.91
##	3rd Qu.:1832	3rd Qu.:1548	3rd Qu.:278.5	3rd Qu.: 72.00
##	Max. :2525	Max. :2170	Max. :376.0	Max. :155.00
##				
##	TEAM_BATTING_HR	TEAM_BATTING_BB	TEAM_BATTING_SO	TEAM_BASERUN_SB
##	Min. : 0.00	Min. : 15.0	Min. : 0.0	Min. : 0.0
##	1st Qu.: 44.50	1st Qu.:436.5	1st Qu.: 545.0	1st Qu.: 59.0
##	Median :101.00	Median :509.0	Median : 686.0	Median : 92.0
##	Mean : 95.63	Mean :499.0	Mean : 709.3	Mean :123.7
##	3rd Qu.:135.50	3rd Qu.:565.5	3rd Qu.: 912.0	3rd Qu.:151.8
##	Max. :242.00	Max. :792.0	Max. :1268.0	Max. :580.0
##			NA's :18	NA's :13
##	TEAM_BASERUN_CS	TEAM_BATTING_HBP	TEAM_PITCHING_H	TEAM_PITCHING_HR
##	Min. : 0.00	Min. :42.00	Min. : 1155	Min. : 0.0
##	1st Qu.: 38.00	1st Qu.:53.50	1st Qu.: 1426	1st Qu.: 52.0
##	Median : 49.50	Median :62.00	Median : 1515	Median :104.0
##	Mean : 52.32	Mean :62.37	Mean : 1813	Mean :102.1
##	3rd Qu.: 63.00	3rd Qu.:67.50	3rd Qu.: 1681	3rd Qu.:142.5
##	Max. :154.00	Max. :96.00	Max. :22768	Max. :336.0
##	NA's :87	NA's :240		
##	TEAM_PITCHING_BB	TEAM_PITCHING_SO	TEAM_FIELDING_E	TEAM_FIELDING_DP
##	Min. : 136.0	Min. : 0.0	Min. : 73.0	Min. : 69.0
##	1st Qu.: 471.0	1st Qu.: 613.0	1st Qu.: 131.0	1st Qu.:131.0
##	Median : 526.0	Median : 745.0	Median : 163.0	Median :148.0
##	Mean : 552.4	Mean : 799.7	Mean : 249.7	Mean :146.1
##	3rd Qu.: 606.5	3rd Qu.: 938.0	3rd Qu.: 252.0	3rd Qu.:164.0
##	Max. :2008.0	Max. :9963.0	Max. :1568.0	Max. :204.0
##		NA's :18		NA's :31

Glimpse of the train data

```
## Rows: 2,276
## Columns: 17
## $ INDEX      <int> 1, 2, 3, 4, 5, 6, 7, 8, 11, 12, 13, 15, 16, 17, 1
## $ TARGET_WINS <int> 39, 70, 86, 70, 82, 75, 80, 85, 86, 76, 78, 68, 7
## $ TEAM_BATTING_H <int> 1445, 1339, 1377, 1387, 1297, 1279, 1244, 1273, 1
## $ TEAM_BATTING_2B <int> 194, 219, 232, 209, 186, 200, 179, 171, 197, 213,
## $ TEAM_BATTING_3B <int> 39, 22, 35, 38, 27, 36, 54, 37, 40, 18, 27, 31, 4
## $ TEAM_BATTING_HR <int> 13, 190, 137, 96, 102, 92, 122, 115, 114, 96, 82,
## $ TEAM_BATTING_BB <int> 143, 685, 602, 451, 472, 443, 525, 456, 447, 441,
## $ TEAM_BATTING_SO <int> 842, 1075, 917, 922, 920, 973, 1062, 1027, 922, 8
## $ TEAM_BASERUN_SB <int> NA, 37, 46, 43, 49, 107, 80, 40, 69, 72, 60, 119,
## $ TEAM_BASERUN_CS <int> NA, 28, 27, 30, 39, 59, 54, 36, 27, 34, 39, 79, 1
## $ TEAM_BATTING_HBP <int> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N
## $ TEAM_PITCHING_H <int> 9364, 1347, 1377, 1396, 1297, 1279, 1244, 1281, 1
## $ TEAM_PITCHING_HR <int> 84, 191, 137, 97, 102, 92, 122, 116, 114, 96, 86,
## $ TEAM_PITCHING_BB <int> 927, 689, 602, 454, 472, 443, 525, 459, 447, 441,
## $ TEAM_PITCHING_SO <int> 5456, 1082, 917, 928, 920, 973, 1062, 1033, 922,
## $ TEAM_FIELDING_E <int> 1011, 193, 175, 164, 138, 123, 136, 112, 127, 131
## $ TEAM_FIELDING_DP <int> NA, 155, 153, 156, 168, 149, 186, 136, 169, 159,
```

Glimpse of the test data

```
## Rows: 259
## Columns: 16
## $ INDEX          <int> 9, 10, 14, 47, 60, 63, 74, 83, 98, 120, 123, 135,
138...
## $ TEAM_BATTING_H  <int> 1209, 1221, 1395, 1539, 1445, 1431, 1430, 1385, 1
259,...
## $ TEAM_BATTING_2B <int> 170, 151, 183, 309, 203, 236, 219, 158, 177, 212,
243...
## $ TEAM_BATTING_3B <int> 33, 29, 29, 29, 68, 53, 55, 42, 78, 42, 40, 55, 5
7, 2...
## $ TEAM_BATTING_HR <int> 83, 88, 93, 159, 5, 10, 37, 33, 23, 58, 50, 164,
186,...
## $ TEAM_BATTING_BB <int> 447, 516, 509, 486, 95, 215, 568, 356, 466, 452,
495,...
## $ TEAM_BATTING_SO <int> 1080, 929, 816, 914, 416, 377, 527, 609, 689, 584
, 64...
## $ TEAM_BASERUN_SB <int> 62, 54, 59, 148, NA, NA, 365, 185, 150, 52, 64, 4
8, 3...
## $ TEAM_BASERUN_CS <int> 50, 39, 47, 57, NA, NA, NA, NA, NA, NA, NA, 28, 2
1, 8...
## $ TEAM_BATTING_HBP <int> NA, NA, NA, 42, NA, NA, NA, NA, NA, NA, NA, NA, N
A, N...
## $ TEAM_PITCHING_H <int> 1209, 1221, 1395, 1539, 3902, 2793, 1544, 1626, 1
342,...
## $ TEAM_PITCHING_HR <int> 83, 88, 93, 159, 14, 20, 40, 39, 25, 62, 53, 173,
196...
## $ TEAM_PITCHING_BB <int> 447, 516, 509, 486, 257, 420, 613, 418, 497, 482,
521...
## $ TEAM_PITCHING_SO <int> 1080, 929, 816, 914, 1123, 736, 569, 715, 734, 62
2, 6...
## $ TEAM_FIELDING_E <int> 140, 135, 156, 124, 616, 572, 490, 328, 226, 184,
200...
## $ TEAM_FIELDING_DP <int> 156, 164, 153, 154, 130, 105, NA, 104, 132, 145,
183,...
```


Find SD for all of the train data

##	INDEX	TEAM_BATTING_H	TEAM_BATTING_2B	TEAM_BATTING_3B
##	693.28867	150.65523	49.51612	27.14410
##	TEAM_BATTING_HR	TEAM_BATTING_BB	TEAM_BATTING_SO	TEAM_BASERUN_SB
##	56.33221	120.59215	243.11114	93.38796
##	TEAM_BASERUN_CS	TEAM_BATTING_HBP	TEAM_PITCHING_H	TEAM_PITCHING_HR
##	23.10457	12.70700	1662.91308	57.65490
##	TEAM_PITCHING_BB	TEAM_PITCHING_SO	TEAM_FIELDING_E	TEAM_FIELDING_DP
##	172.95006	634.30585	230.90260	25.88387

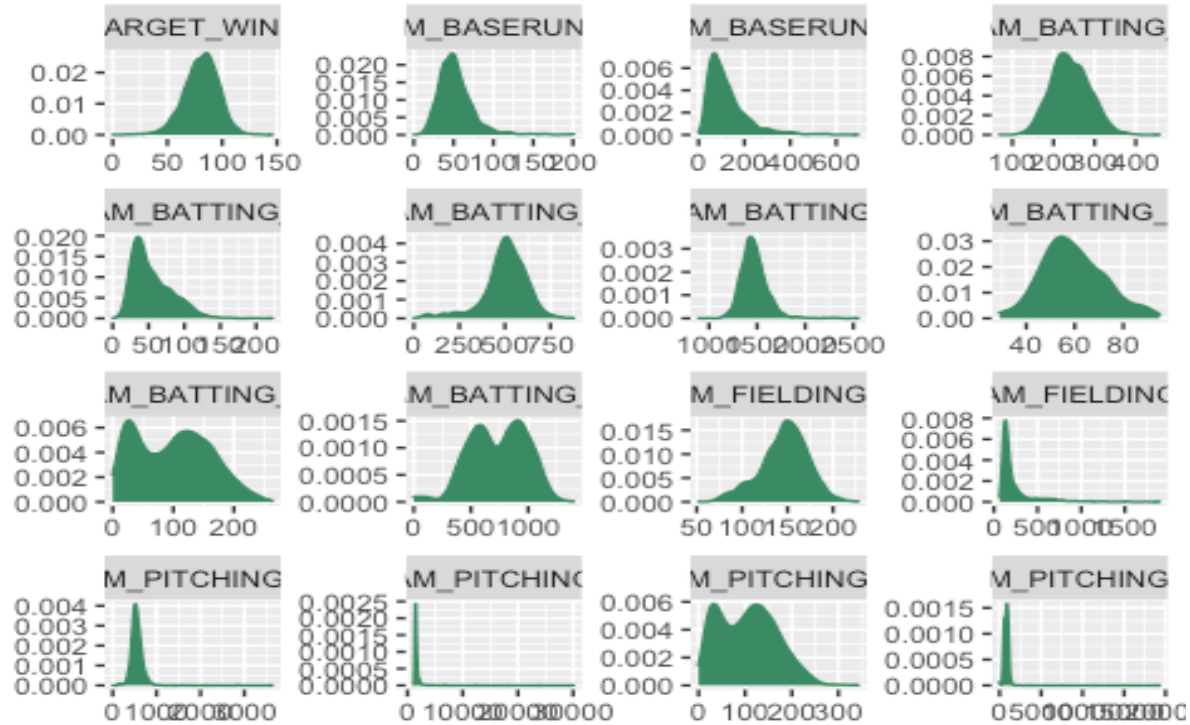
Find SD for all of the test data

##	INDEX	TARGET_WINS	TEAM_BATTING_H	TEAM_BATTING_2B
##	736.34904	15.75215	144.59120	46.80141
##	TEAM_BATTING_3B	TEAM_BATTING_HR	TEAM_BATTING_BB	TEAM_BATTING_SO
##	27.93856	60.54687	122.67086	248.52642
##	TEAM_BASERUN_SB	TEAM_BASERUN_CS	TEAM_BATTING_HBP	TEAM_PITCHING_H
##	87.79117	22.95634	12.96712	1406.84293
##	TEAM_PITCHING_HR	TEAM_PITCHING_BB	TEAM_PITCHING_SO	TEAM_FIELDING_E
##	61.29875	166.35736	553.08503	227.77097
##	TEAM_FIELDING_DP			
##	26.22639			

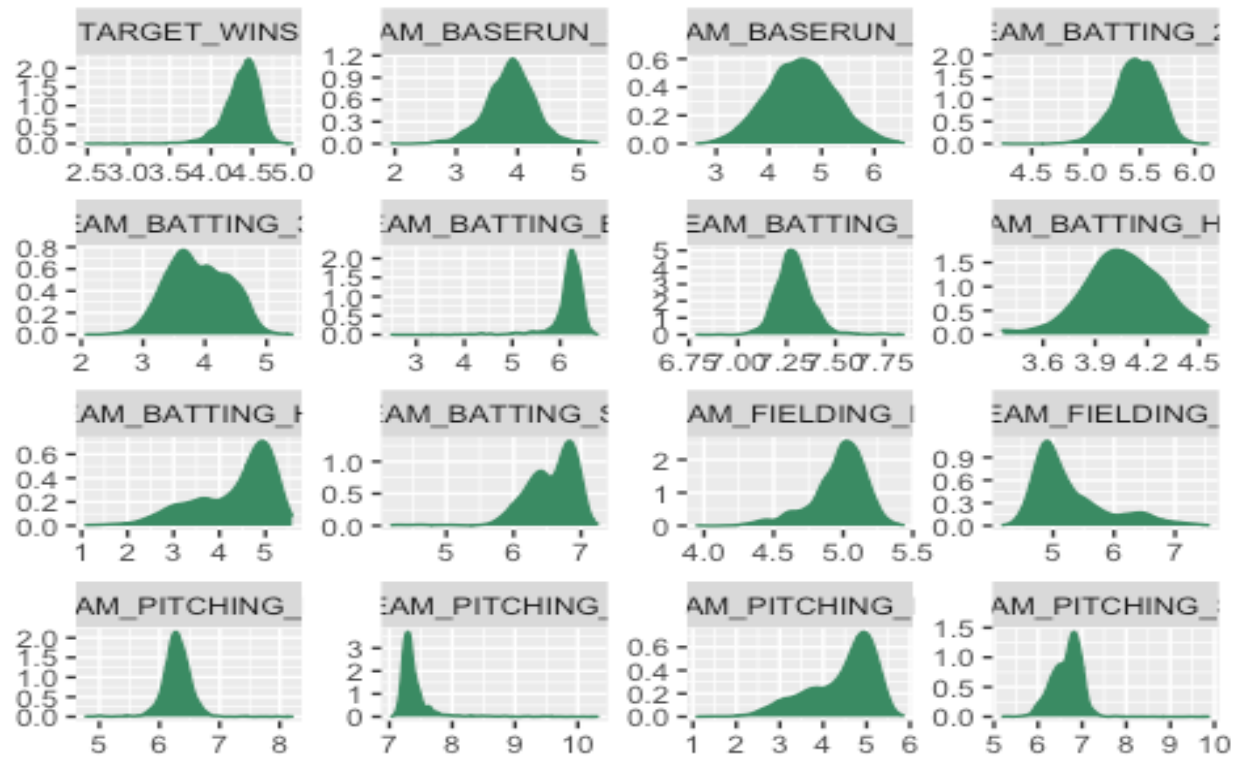
Box plot the train data




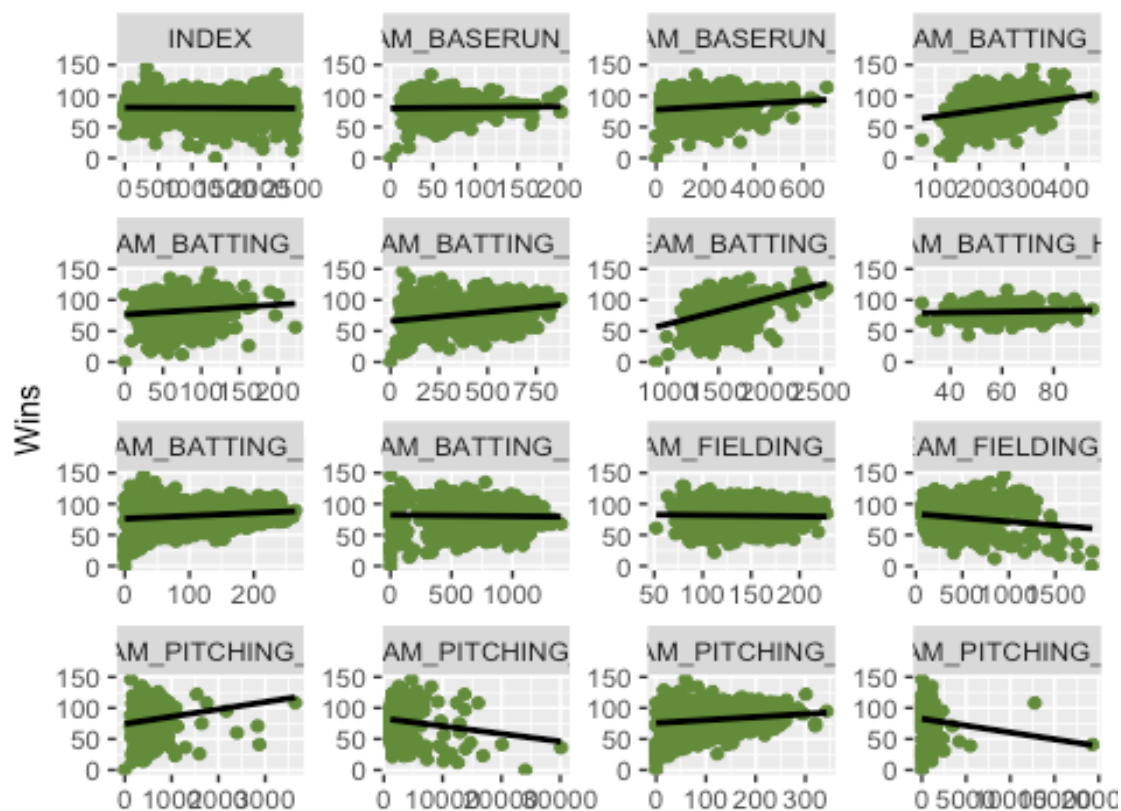
Variable Distributions



Log Variable Distributions



Correlations with Response Variable




DATA PREPARATION

NA counts for the train data set ¹

##	INDEX	TARGET_WINS	TEAM_BATTING_H	TEAM_BATTING_2B
##	0	0	0	0
##	TEAM_BATTING_3B	TEAM_BATTING_HR	TEAM_BATTING_BB	TEAM_BATTING_SO
##	0	0	0	102
##	TEAM_BASERUN_SB	TEAM_BASERUN_CS	TEAM_BATTING_HBP	TEAM_PITCHING_H
##	131	772	2085	0
##	TEAM_PITCHING_HR	TEAM_PITCHING_BB	TEAM_PITCHING_SO	TEAM_FIELDING_E
##	0	0	102	0
##	TEAM_FIELDING_DP			
##	286			


visulaization and percentage of NA values ²



¹ <https://statisticsglobe.com/count-number-of-na-values-in-vector-and-column-in-r>

² <https://cran.r-project.org/web/packages/naniar/vignettes/naniar-visualisation.html>

alternative NA values visualization ³



Since 92% of the data for the TEAM_BATTING_HBP is missing, the variable has been removed from both test and train data. TEAM_BASERUN_CS is a runner up with the next highest amount of NA at 34%.

³ <https://datavizpyr.com/visualizing-missing-data-with-barplot-in-r/>

BUILD MODELS

Model #1

Two predictors: Base hits by batters and Hits allowed

Using a manual review, below are the features selected for the first model and the supporting reason/s.

TEAM_BATTING_H = Base hits by batters: it's impossible to win in baseball without getting to the bases and hitting the ball is the primary means to accomplish this.

TEAM_PITCHING_H = Hits allowed: winning without a good defense is difficult and in baseball preventing the other team from getting hits is a good defense strategy.

Only two features are selected for the first model - start small and build up seems like a good approach.

Create the Regression Model

```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + TEAM_PITCHING_H,
##     data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -57.444  -9.113   0.613   9.677  77.744
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   15.6425241   3.5652886   4.387 1.22e-05 ***
## TEAM_BATTING_H    0.0475405   0.0024789  19.178 < 2e-16 ***
## TEAM_PITCHING_H -0.0026104   0.0002462 -10.603 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 14.16 on 1704 degrees of freedom
## Multiple R-squared:  0.189, Adjusted R-squared:  0.1881
## F-statistic: 198.6 on 2 and 1704 DF, p-value: < 2.2e-16
```

The p values are 0, which per the criteria of “keep a feature if the p-value is <0.05” recommends that we keep both these features. But, the adjusted R-squared is TERRIBLE at around 21%. Even though the R-squared is poor it's simple to run this model with the test data, so we'll do that next.

Evaluate the second model results using RMSE

```
## [1] 13.6336
```

Model #2

Four predictors: Base hits by batters, Hits allowed, Errors, and Walks allowed

Using a manual review, below are the features selected for the second model and the supporting reason/s.

We'll keep the features from the first model (due to low p-values) and add two more features... TEAM_FIELDING_E = Errors: errors are costly in terms of immediate impact, but could also impact the team in other ways (i.e. a high occurrence could impact team comraderie and confidence in each other)

TEAM_PITCHING_BB = Walks allowed: putting players on base for "free" is more opportunity for points

Create the Regression Model

```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + TEAM_PITCHING_H +
##     TEAM_FIELDING_E + TEAM_PITCHING_BB, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -51.927  -9.234   0.174   9.477  47.151
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    7.2608113   3.6224175   2.004  0.0452 *
## TEAM_BATTING_H    0.0495348   0.0024231  20.443 < 2e-16 ***
## TEAM_PITCHING_H  -0.0018244   0.0003488  -5.231 1.89e-07 ***
## TEAM_FIELDING_E  -0.0129111   0.0020960  -6.160 9.07e-10 ***
## TEAM_PITCHING_BB  0.0130541   0.0023301   5.602 2.46e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.73 on 1702 degrees of freedom
## Multiple R-squared:  0.2382, Adjusted R-squared:  0.2364
## F-statistic: 133 on 4 and 1702 DF, p-value: < 2.2e-16
```

Evaluate the second model results using RMSE

```
## [1] 13.30535
```

The increase from two features in the first model to four features in the second model did not yield a noticeable improvement. The Adjusted R2 on the training data improved slightly, but the RMSE for all practical purposes stayed the same at around 13; which is a poor RMSE implying that both models have poor predictive capability.

Model #3

BSR Model (SaberMetrics) (data imputation)

Base runs (BsR) is a baseball statistic invented by sabermetrician David Smyth to estimate the number of runs a team “should have” scored given their component offensive statistics, as well as the number of runs a hitter or pitcher creates or allows. It measures essentially the same thing as Bill James runs created, but as sabermetrician Tom M. Tango points out, base runs models the reality of the run-scoring process “significantly better than any other run estimator”.

Cleaning Data

```
##
## iter imp variable
## 1 1 TEAM_BATTING_SO TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_PITCHING_
SO TEAM_FIELDING_DP
## 2 1 TEAM_BATTING_SO TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_PITCHING_
SO TEAM_FIELDING_DP
## 3 1 TEAM_BATTING_SO TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_PITCHING_
SO TEAM_FIELDING_DP
## 4 1 TEAM_BATTING_SO TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_PITCHING_
SO TEAM_FIELDING_DP
## 5 1 TEAM_BATTING_SO TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_PITCHING_
SO TEAM_FIELDING_DP
```

The simplest, uses only the most common batting statistics[2]

$$A = H + BB - HR \quad B = (1.4 * TB - .6 * H - 3 * HR + .1 * BB) * 1.02 \quad C = AB - H \quad D = HR$$

$$BsR = \frac{(A * B)}{(B + C)} + D$$

Create the Regression Model

BSR

```
##
## Call:
## lm(formula = TARGET_WINS ~ BSR + TEAM_PITCHING_SO + TEAM_FIELDING_E +
##     TEAM_FIELDING_DP, data = rmdata3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -63.776  -8.418   0.410   8.537  50.024
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   45.237755   3.161743  14.308 < 2e-16 ***
## BSR           0.049301   0.002015  24.469 < 2e-16 ***
## TEAM_PITCHING_SO 0.009641   0.001308   7.369 2.67e-13 ***
## TEAM_FIELDING_E -0.039350   0.001995 -19.728 < 2e-16 ***
```

```
## TEAM_FIELDING_DP -0.169143    0.012539 -13.489  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.18 on 1702 degrees of freedom
## Multiple R-squared:  0.3157, Adjusted R-squared:  0.3141
## F-statistic: 196.3 on 4 and 1702 DF,  p-value: < 2.2e-16
```

Evaluate the model results using RMSE

```
## [1] 14.25345
```

Model #4

(Modified) Backward Elimination Model (omitting NAs)

Due to previously learning how to perform Backward Elimination and it being possible to perform manually, we decided to include a model that resulted from the procedure. The process was performed with imputed data (via MICE) as well as data with NAs removed. The latter showed stronger results, therefore the final model was fitted with the NA omitted data.

According to Faraway, Backward Elimination is when you start with all predictors in the model, then remove the predictor with the highest p-value as long as it is above your p-value threshold (e.g. 0.05). Then refit the model and continue the process until only predictors with p-values below your threshold remain.

Additionally, we took steps to remove variables with non-intuitive coefficients. For instance, TEAM_FIELDING_DP and TEAM_PITCHING_SO were unexpectedly showing negative effects on wins. While there could be potential intervening variables giving these variables true predictive power, we opted to remove the variables from the model due to the possibility they were significant by chance and due to our bias towards parsimony. Further, RMSE did not drastically worsen when removed.


```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BASERUN_SB + TEAM_BATTING_HR +
##     TEAM_BATTING_BB + TEAM_FIELDING_E, data = test_no_na)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -38.055  -7.316   0.722   6.485  26.008
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   56.249572   6.027751   9.332  < 2e-16 ***
## TEAM_BASERUN_SB  0.047422   0.013540   3.502 0.000520 ***
## TEAM_BATTING_HR  0.055258   0.016007   3.452 0.000622 ***
## TEAM_BATTING_BB  0.035161   0.007576   4.641 4.87e-06 ***
## TEAM_FIELDING_E -0.045877   0.018742  -2.448 0.014853 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.16 on 359 degrees of freedom
## Multiple R-squared:  0.2096, Adjusted R-squared:  0.2008
## F-statistic: 23.8 on 4 and 359 DF,  p-value: < 2.2e-16
##
## [1] 11.081
```


SELECT MODELS

Verifying OLS Regression Assumptions



Assumption: Mean of residuals is zero


```
## [1] 2.396206e-17
```



Fitted values


INS ~ TEAM_BASERUN_SB + TEAM_BATTING_HR + TEAM_BAT





INS ~ TEAM_BASERUN_SB + TEAM_BATTING_HR + TEAM_BAT

Series residuals(backward_mod_model)



Model Selection

First, before fully evaluating models we validated that all individual predictors had p-values below 0.05, the cutoff for a 95% confidence level. Additionally, we validated that the models F-statistics were also significant at a 95% confidence level.

Then, the two primary statistics used to choose our final model were adjusted R-squared and root mean square error (RMSE). Adjusted R-squared helped guide model selection since, like R-squared, adjusted R-squared measures the amount of variation in the dependent variable explained by the independent variables, except with a correction to ensure only independent variables with predictive power raise the statistic. RMSE was perhaps even more crucial to model selection as it is the measure of the standard deviation of the residuals, essentially a measure of accuracy in the same units as the response variable. To ensure the model can generalize to unobserved data, we calculated the RMSE on our test set.

Backward elimination saw a RMSE of approximately 10, noticeably outperforming other models. Therefore, we chose the backward elimination model even with a slightly worse adjusted R-squared. Additionally, since all top performing models included four predictors, parsimony was not a consideration.

Lastly, we verified the forward selection model meets OLS regression assumptions. These included: no significant multicollinearity, the mean of residuals is zero, homoscedasticity of residuals, and no significant auto-correlation. We deemed all assumptions had been met, but note, there is a slight trend in the residuals vs fitted plot (Assumption: Homoscedasticity of residuals) which may indicate a small nonlinear trend.

Matt's References

Bhandari, Aniruddha, "Key Difference between R-squared and Adjusted R-squared for Regression Analysis", Analytics Vidhya, 2020
<https://www.analyticsvidhya.com/blog/2020/07/difference-between-r-squared-and-adjusted-r-squared/>

Glen., Stephanie "RMSE: Root Mean Square Error", StatisticsHowTo.com
<https://www.statisticshowto.com/probability-and-statistics/regression-analysis/rmse-root-mean-square-error/>

Gupta, Aryansh, "Linear Regression Assumptions and Diagnostics in R", RPubS,
<https://rpubs.com/aryn999/LinearRegressionAssumptionsAndDiagnosticsInR>

Kim, Bommae, "Understanding Diagnostic Plots for Linear Regression Analysis", University of Virginia Library, <https://data.library.virginia.edu/diagnostic-plots/>

Code Appendix

DATA EXPLORATION:

#We can observe the response variable (TARGET_WINS) Looks to be normally distributed. This supports the working theory that there are good teams and bad teams. There are also a lot of average teams.

#There are also quite a few variables with missing values. and, Some variables are right skewed (TEAM_BASERUN_CS, TEAM_BASERUN_SB, etc.). This might support the good team theory. It may also introduce non-normally distributed residuals in the model. We shall see.

Load the Data

Set seed for reproducibility
`set.seed(621)`

`train <- read.csv("https://raw.githubusercontent.com/akarimhammoud/Data_621/main/Assignment_1/data/moneyball-training-data.csv")`

`evaluation <- read.csv("https://raw.githubusercontent.com/akarimhammoud/Data_621/main/Assignment_1/data/moneyball-evaluation-data.csv")`

Summary of the data

`summary(train)`
`summary(evaluation)`

Glimpse of the data

`glimpse(train)`

`glimpse(evaluation)`

Find SD for all of the train and test data

`apply(train, 2, sd, na.rm=TRUE)`

`apply(evaluation, 2, sd, na.rm=TRUE)`

Box plot the data

```
ggplot(stack(train), aes(x = ind, y = values)) +  
  geom_boxplot() +  
  theme(legend.position="none") +  
  theme(axis.text.x=element_text(angle=45, hjust=1))
```

Variable Distributions

```
train %>%  
  gather(variable, value, TARGET_WINS:TEAM_FIELDING_DP) %>%  
  ggplot(., aes(value)) +  
  geom_density(fill = "#3A8B63", color="#3A8B63") +  
  facet_wrap(~variable, scales = "free", ncol = 4) +  
  labs(x = element_blank(), y = element_blank())
```

#Log Variable Distributions

```
train_log <- log(train)  
  
train_log %>%  
  gather(variable, value, TARGET_WINS:TEAM_FIELDING_DP) %>%  
  ggplot(., aes(value)) +  
  geom_density(fill = "#3A8B63", color="#3A8B63") +  
  facet_wrap(~variable, scales = "free", ncol = 4) +  
  labs(x = element_blank(), y = element_blank())
```

Correlations with Response Variable

```
train %>%  
  gather(variable, value, -TARGET_WINS) %>%  
  ggplot(., aes(value, TARGET_WINS)) +  
  geom_point(fill = "#628B3A", color="#628B3A") +  
  geom_smooth(method = "lm", se = FALSE, color = "black") +  
  facet_wrap(~variable, scales = "free", ncol = 4) +  
  labs(x = element_blank(), y = "Wins")  
  
train %>%  
  cor(., use = "complete.obs") %>%  
  corplot(., method = "color", type = "upper", tl.col = "black", diag = FALSE)
```

DATA PREPARATION

```
# ^[https://statisticsglobe.com/count-number-of-na-values-in-vector-and-column-in-r/]
```

```
#NA counts for the train data set  
colSums(is.na(train))
```

```
# ^[https://cran.r-project.org/web/packages/naniar/vignettes/naniar-visualisation.html]
```

```
#visulaization and percentage of NA values  
vis_miss(train)
```

```
# ^[https://datavizpyr.com/visualizing-missing-data-with-barplot-in-r/]
```

```
#alternative NA values visualization
```

```
train %>%  
  summarise_all(list(~is.na(.)))%>%  
  pivot_longer(everything(),  
               names_to = "variables", values_to="missing") %>%  
  count(variables, missing) %>%  
  ggplot(aes(y=variables,x=n,fill=missing))+  
  geom_col()
```

#Since 92% of the data for the TEAM_BATTING_HBP is missing, the variable has been removed from both test #and train data. TEAM_BASERUN_CS is a runner up w ith the next highest amount of NA at 34%.

```
#removes the TEAM_BATTING_HBP due to high # of NAs
```

```
train_full <- train %>% dplyr::select(-c(Team_BATTING_HBP))  
evaluation <- evaluation %>% dplyr::select(-c(Team_BATTING_HBP))
```

```
# ^[https://sphweb.bumc.bu.edu/otlt/MPH-Modules/BS/R/R-Manual/R-Manual5.html]
```

```
#creates CSV in your current working directory of R
```

```
write.csv(train_full, 'hw1_train_data.csv')  
write.csv(evaluation, 'hw1_evaluation_data.csv')
```

```
# Create train, test split
```

```
train <- train_full %>% dplyr::sample_frac(.75)
test  <- dplyr::anti_join(train_full, train, by = 'INDEX')
```

BUILD MODELS

Model #1

Two predictors: Base hits by batters and Hits allowed

#Using a manual review, below are the features selected for the first model and the supporting reason/s.

#TEAM_BATTING_H = Base hits by batters: it's impossible to win in baseball without getting to the bases # and hitting the ball is the primary means to accomplish this.

#TEAM_PITCHING_H = Hits allowed: winning without a good defense is difficult and in baseball preventing the other team from getting hits is a good defensive strategy.

#Only two features are selected for the first model - start small and build up seems like a good approach.

*# Create the Regression Model *

Build the first model and produce a summary

```
first_model <- lm(TARGET_WINS ~ TEAM_BATTING_H + TEAM_PITCHING_H, data = train)
summary(first_model)
```

#The p values are 0, which per the criteria of "keep a feature if the p-value is <0.05" recommends that we keep both these features. But, the adjusted R-squared is TERRIBLE at around 21%. Even though the R-squared is poor it's simple to run this model with the test data, so we'll do that next.

#Predict with the first model training data

```
first_model_predictions = predict(first_model, test)
```

#Evaluate the first model results using RMSE

```
rmse(test$TARGET_WINS, first_model_predictions)
```

Model #2

Four predictors: Base hits by batters, Hits allowed, Errors, and Walks allowed

#Using a manual review, below are the features selected for the second model and the supporting reason/s.

#We'll keep the features from the first model (due to low p-values) and add two more features...

#TEAM_FIELDING_E = Errors: errors are costly in terms of immediate impact, but could also impact the team in other ways (i.e. a high occurrence could impact team comradery and confidence in each other)

#TEAM_PITCHING_BB = Walks allowed: putting players on base for "free" is more opportunity for points

Create the Regression Model

Build the second model and produce a summary

```
second_model <- lm(TARGET_WINS ~ TEAM_BATTING_H + TEAM_PITCHING_H + TEAM_FIELDING_E + TEAM_PITCHING_BB, data = train)
summary(second_model)
```

#Predict with the second model training data

```
second_model_predictions = predict(second_model, test)
```

#Evaluate the second model results using RMSE

```
rmse(test$TARGET_WINS, second_model_predictions)
```

#The increase from two features in the first model to four features in the second model did not yield a noticeable improvement. The Adjusted R² on the training data improved slightly, but the RMSE for all practical purposes stayed the same at around 13; which is a poor RMSE implying that both models have poor predictive capability.

Model #3

BSR Model (SaberMetrics) (data imputation)

*Base runs (BsR) is a baseball statistic invented by sabermetrician David Smyth to estimate the number of runs a team "should have"

#*scored given their component offensive statistics, as well as the number of runs a hitter or pitcher creates or allows.*

#*It measures essentially the same thing as Bill James runs created, but as sabermetrician Tom M. Tango points out, base*

#*runs models the reality of the run-scoring process "significantly better than any other run estimator".*

#*Cleaning Data*

Load data

```

data <- read.csv('hw1_train_data.csv')

#imput data by regression:
data_imp <- mice(data, method = "norm.predict", m = 1)

#complete data
data_complete <- complete(data_imp)

# The simplest, uses only the most common batting statistics[2]

#$A = H + BB - HR$
#$B = (1.4 * TB - .6 * H - 3 * HR + .1 * BB) * 1.02$
#$C = AB - H$
#$D = HR$

#$BSR = \frac{(A * B)}{(B + C)} + D$

data3 <- data_complete %>%
  rowwise() %>%
  mutate(TEAM_BATTING_AB = sum( TEAM_BATTING_H,TEAM_BATTING_BB,TEAM_BATTING_S
O, na.rm=TRUE),
         TEAM_BATTING_1B = TEAM_BATTING_H - (TEAM_BATTING_2B + TEAM_BATTING_3
B + TEAM_BATTING_HR),
         TEAM_BATTING_TB = TEAM_BATTING_1B + (2 * TEAM_BATTING_2B) + (3 * TEA
M_BATTING_3B) + (4 * TEAM_BATTING_HR),
         BSR_A = TEAM_BATTING_H + TEAM_BATTING_BB - TEAM_BATTING_HR,
         BSR_B = (( 1.4 * TEAM_BATTING_TB) - ( 0.6 * TEAM_BATTING_H) - (3 * T
EAM_BATTING_HR) + (0.1 * TEAM_BATTING_BB)) * 1.02,
         BSR_C = TEAM_BATTING_AB - TEAM_BATTING_H,
         BSR = ((BSR_A*BSR_B)/(BSR_B + BSR_C)) + TEAM_BATTING_HR
        )

data3 <- as.data.frame(data3)
train3 <- data3 %>% dplyr::sample_frac(.75)
test3  <- dplyr::anti_join(data3, train3, by = 'X')

#<B> Create the Regression Model </B>
#*BSR*

rmdata3 <- train3 %>%
  dplyr::select(BSR, TEAM_PITCHING_SO, TEAM_FIELDING_E, TEAM_FIELDING_DP, TAR
GET_WINS)

#Build the second model and produce a summary
GModel3 <- lm(TARGET_WINS ~ BSR + TEAM_PITCHING_SO + TEAM_FIELDING_E + TEAM_F
IELDING_DP, data = rmdata3)

```

```
summary(GModel3)
```

```
#Predict with the second model training data
```

```
GModel3_predictions = predict(GModel3,test3)
```

```
#Evaluate the second model results using RMSE
```

```
rmse(test3$TARGET_WINS, GModel3_predictions)
```

```
## Model #4
```

```
### (Modified) Backward Elimination Model (omitting NAs)
```

#Due to previously learning how to perform Backward Elimination and it being possible to perform manually, we decided to include a model that resulted from the procedure. The process was performed with imputed data (via MICE) as well as data with NAs removed. The latter showed stronger results, therefore the final model was fitted with the NA omitted data.

#According to Faraway, Backward Elimination is when you start with all predictors in the model, then remove the predictor with the highest p-value as long as it is above your p-value threshold (e.g. 0.05). Then refit the model and continue the process until only predictors with p-values below your threshold remain.

#Additionally, we took steps to remove variables with non-intuitive coefficients. For instance, TEAM_FIELDING_DP and TEAM_PITCHING_SO were unexpectedly showing negative effects on wins. While there could be potential intervening variables giving these variables true predictive power, we opted to remove the variables from the model due to the possibility they were significant by chance and due to our bias towards parsimony. Further, RMSE did not drastically worsen when removed.

```
# Remove NAs
```

```
train_no_na <- na.omit(train)
```

```
test_no_na <- na.omit(test)
```

```
# Fit model
```

```
backward_model <- lm(TARGET_WINS ~ TEAM_BASERUN_SB + TEAM_BATTING_HR + TEAM_BATTING_BB + TEAM_BASERUN_SB  
                      + TEAM_PITCHING_SO + TEAM_FIELDING_E + TEAM_FIELDING_DP,  
data = test_no_na)
```

```
# Fit modified model
```

```

backward_mod_model <- lm(TARGET_WINS ~ TEAM_BASERUN_SB + TEAM_BATTING_HR + TE
AM_BATTING_BB + TEAM_FIELDING_E,
                        data = test_no_na)

# View summary
summary(backward_mod_model)

# Make predictions on test set
backward_model_predictions = predict(backward_mod_model, test_no_na)

# Obtain RMSE between actuals and predicted
rmse(test_no_na$TARGET_WINS, backward_model_predictions)

# Make predictions on evaluation data
backward_model_predictions_evaluation = predict(backward_mod_model, evaluatio
n)

# Final predictions on evaluation set
write.csv(backward_model_predictions_evaluation, 'evaluation_predictions.csv'
)

## SELECT MODELS

### Verifying OLS Regression Assumptions

# Assumption: No Multicollinearity (VIF under 5)
vif(backward_mod_model)

# Assumption: Mean of residuals is zero
mean(residuals(backward_mod_model))

# Assumption: Homoscedasticity of residuals
plot(backward_mod_model)

# Assumption: No auto-correlation
acf(residuals(backward_mod_model), lags=20)

```